

# AXA Data Challenge

Yellow – Zhengying Liu, Chia-Man Hung, Maxime Bellec

# Summary

- Data pre-processing & Feature engineering
- Data visualization & Preliminary analysis
- Our approaches:
  - A first simple approach
  - A generalized version: Linear LinEx Regression
  - Random Forest
- Conclusion

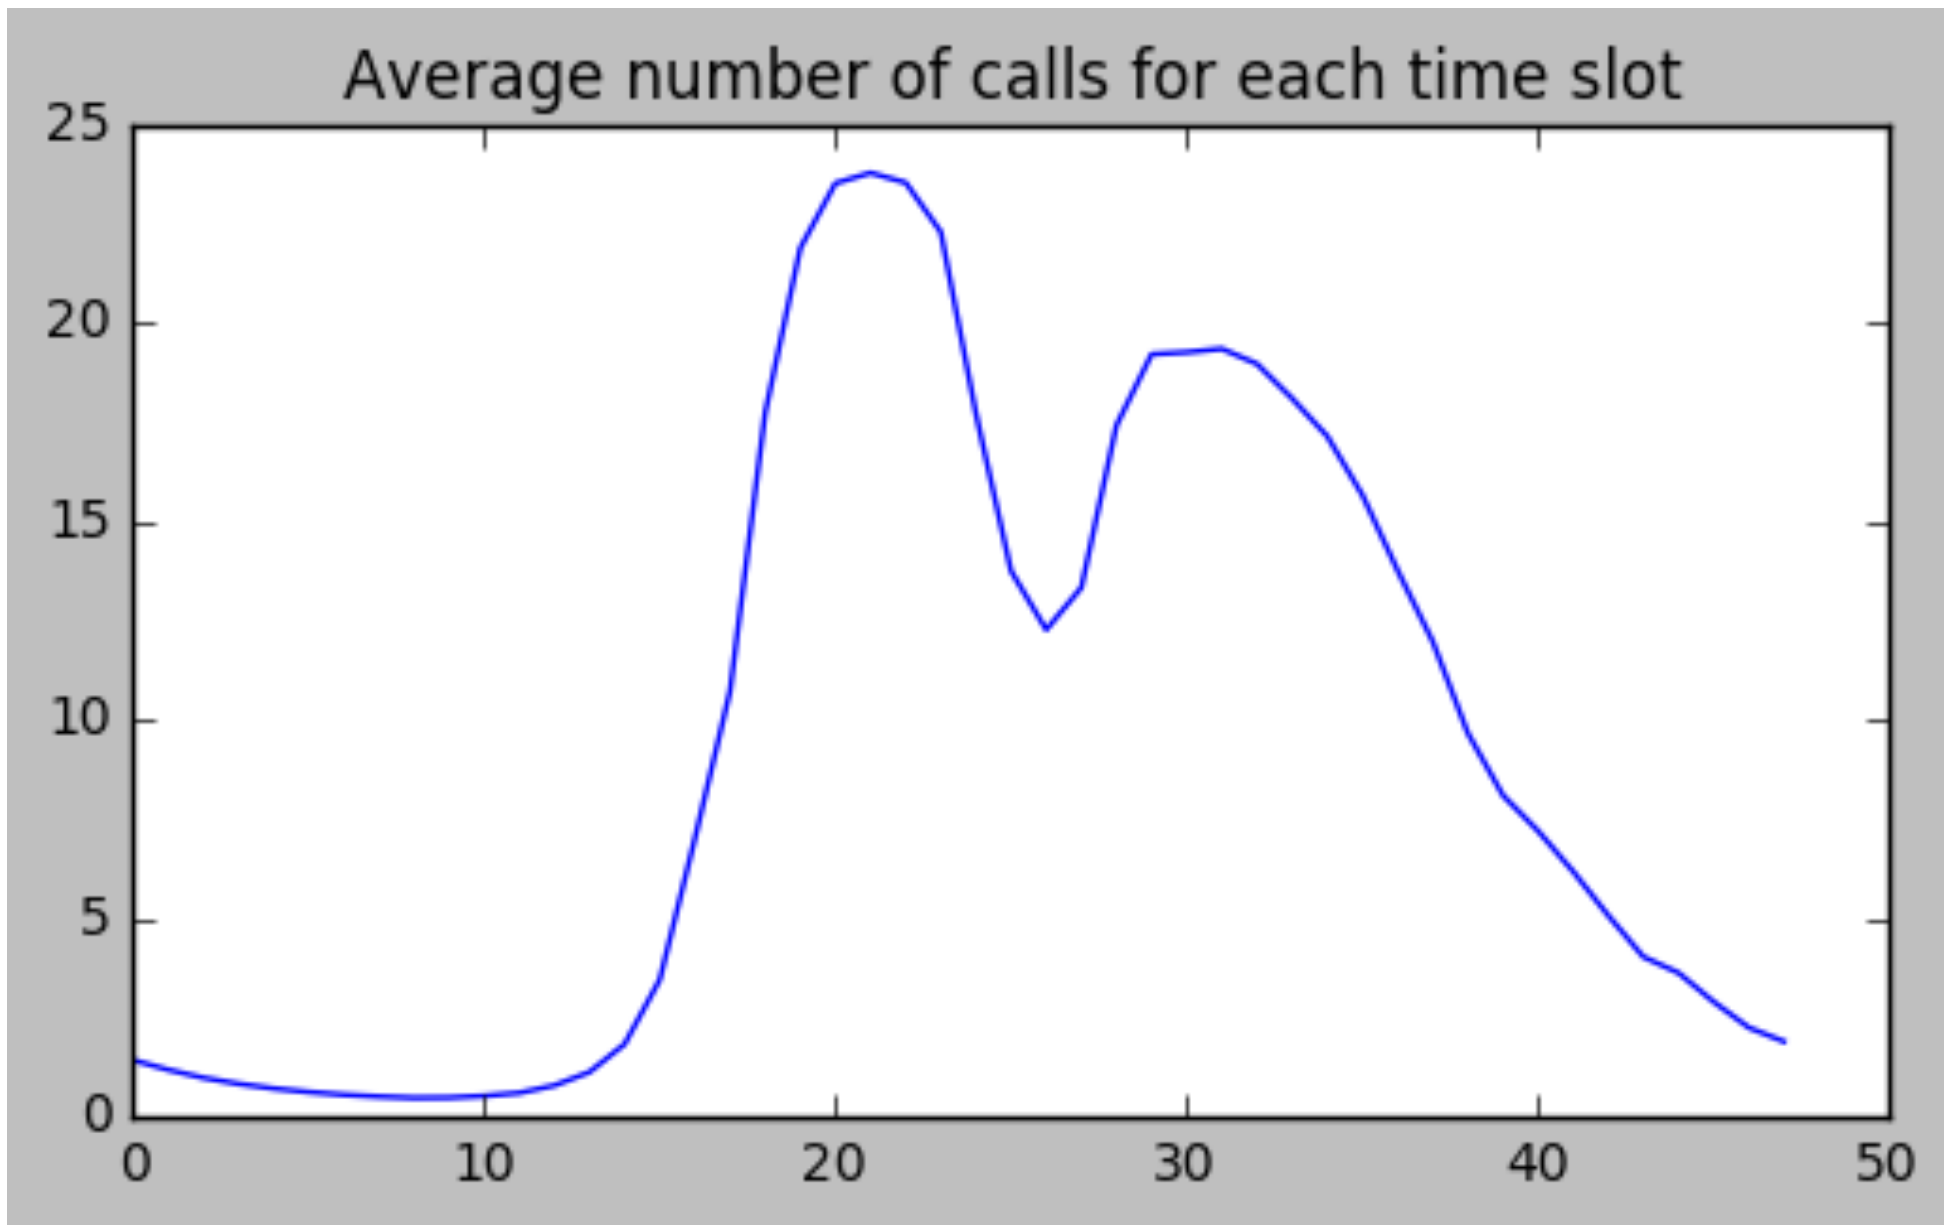
# Data pre-processing

- All columns of the training data are not in test data  
=> they are not useful, except for DATE
- Multiple columns for each (date, ASS\_ASSIGNMENT)  
=> sum the DSLP\_RECEIVED\_CALLS


# Feature engineering

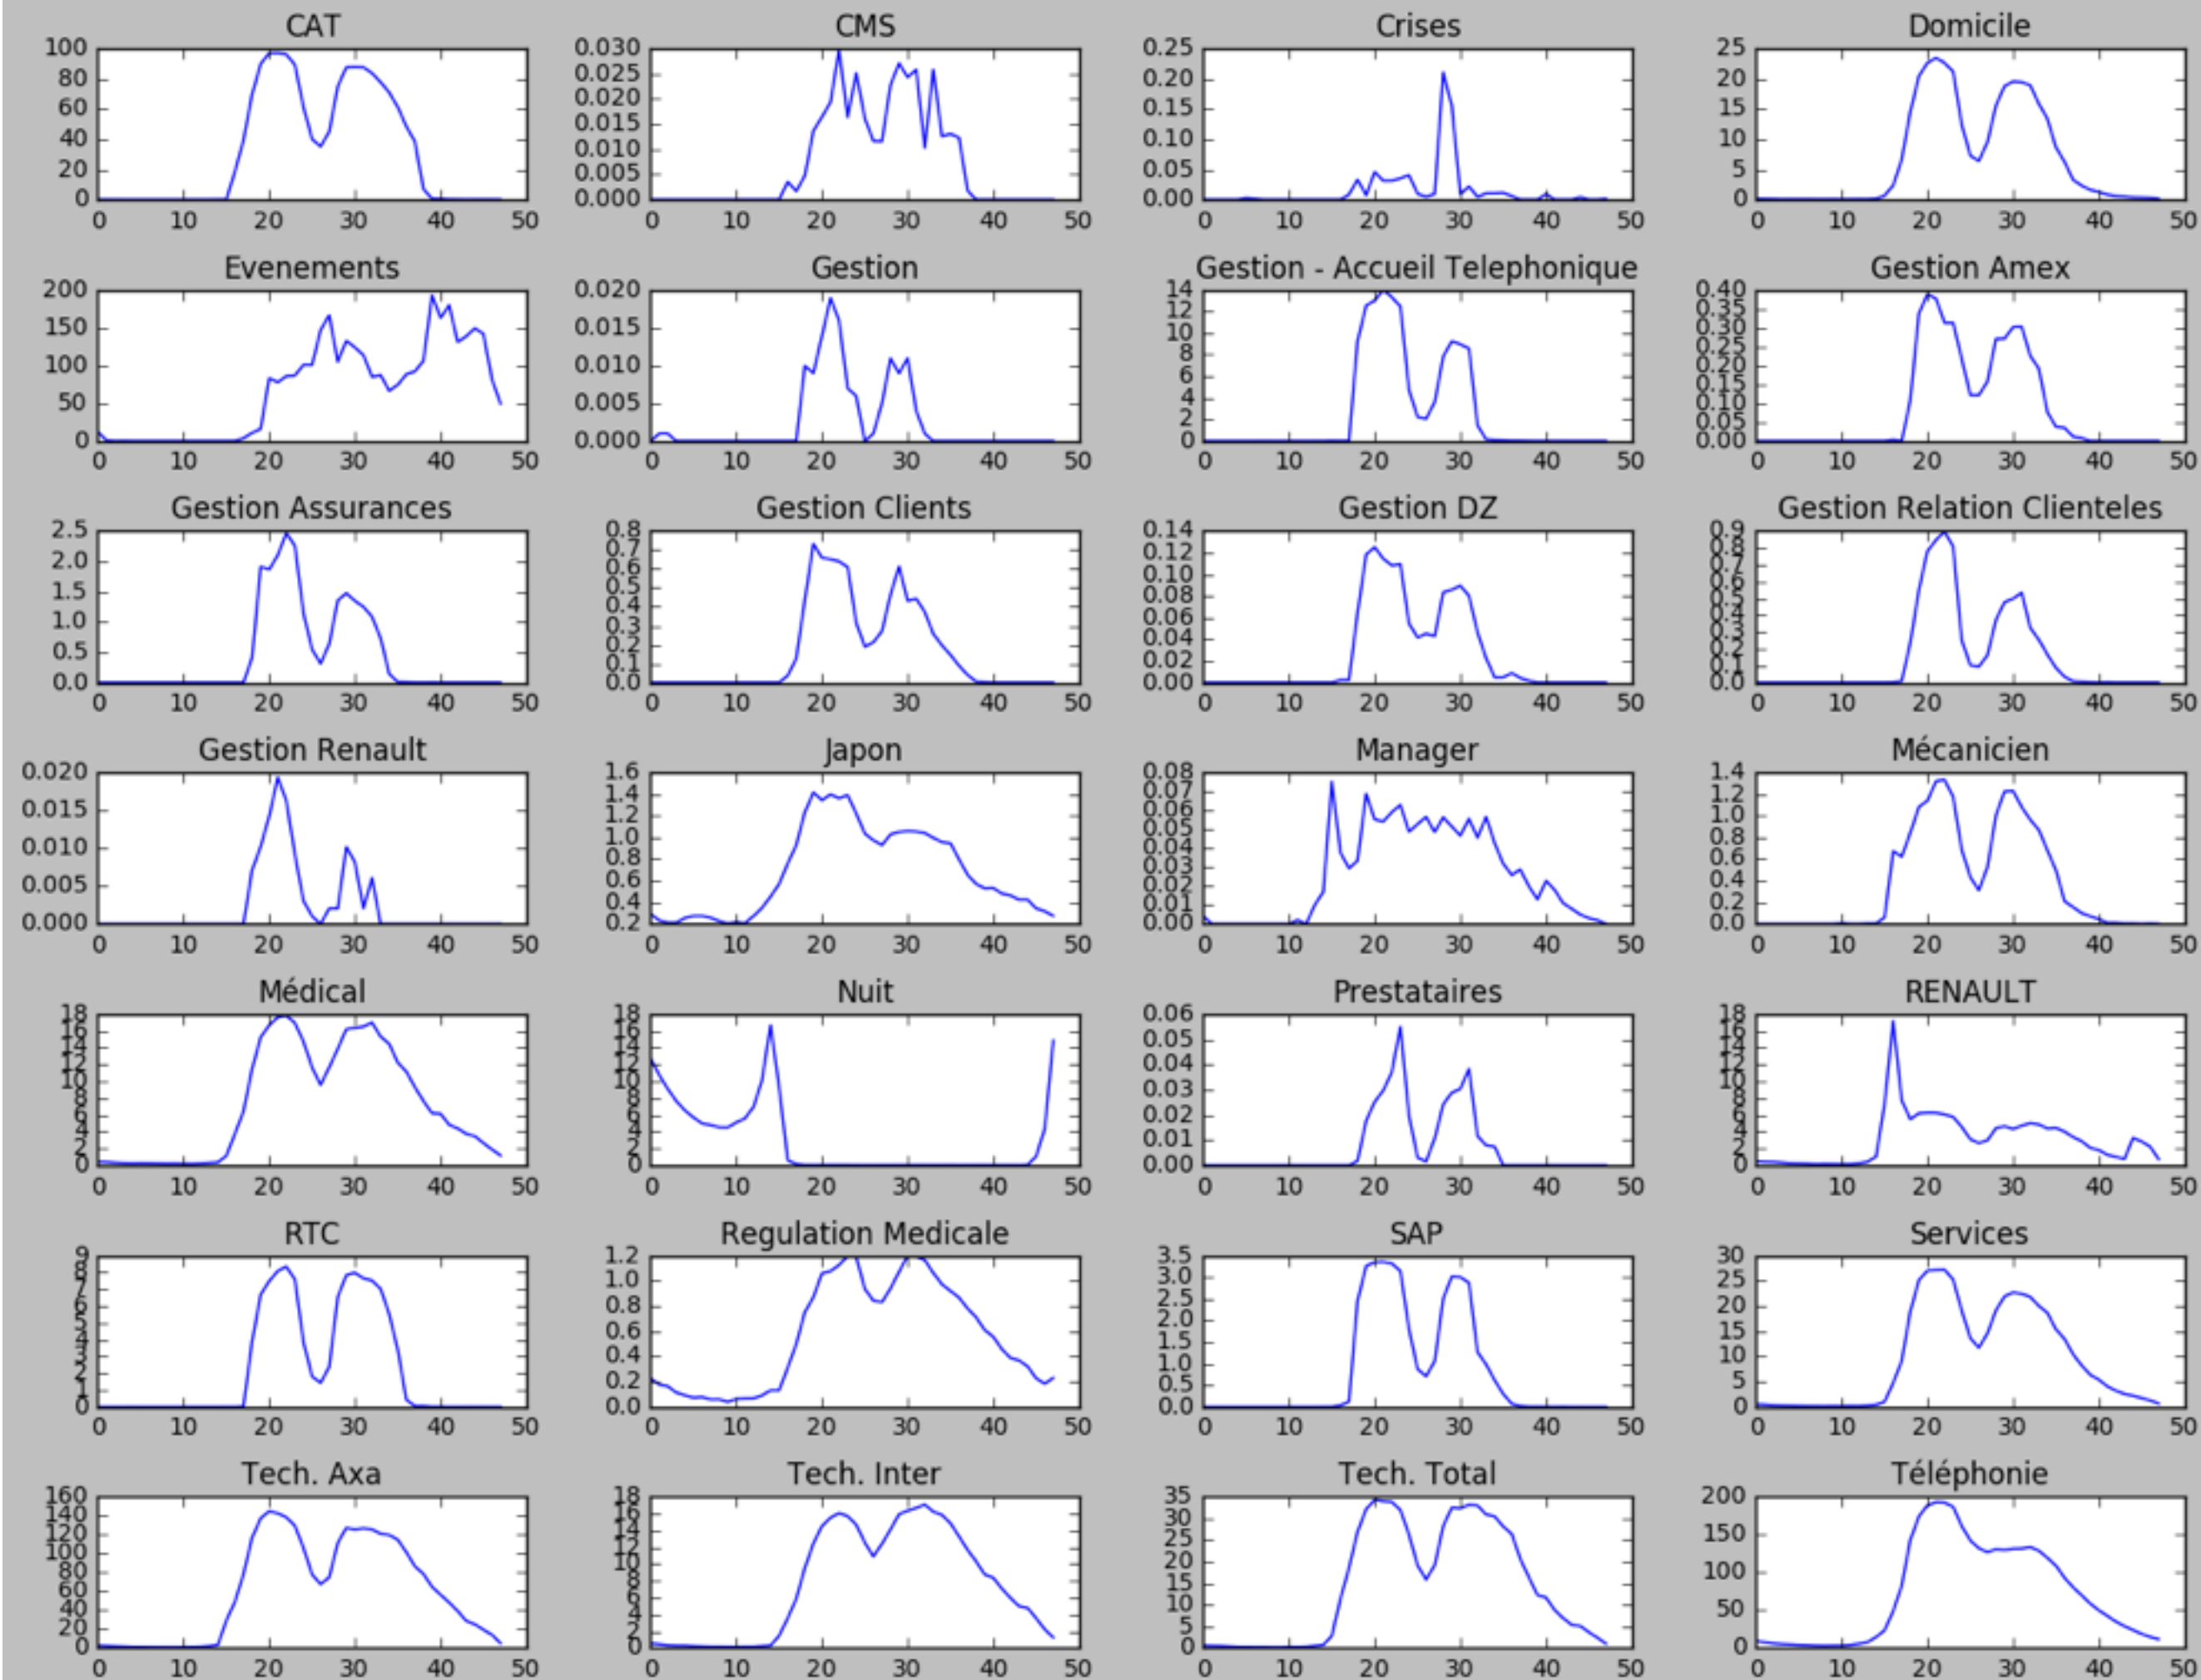
	DATE	ASS_ASSIGNMENT	CSPL_RECEIVED_CALLS	slot	dayofweek	month	year	day_off	day_after_day_off
0	2011-01-01	Crises	0	0	5	1	2011	True	False
1	2011-01-01	Domicile	0	0	5	1	2011	True	False
2	2011-01-01	Gestion	0	0	5	1	2011	True	False
3	2011-01-01	Gestion - Accueil Telephonique	0	0	5	1	2011	True	False
4	2011-01-01	Gestion Amex	0	0	5	1	2011	True	False

# Data visualization



# Data visualization





# 1<sup>st</sup> Approach – simple approach

- For a given ASS\_ASSIGNMENT and weekly time slot, such as Tuesday 09:00-09:30, the CSPL\_RECEIVED\_CALLS are +- stationary
- The idea is to predict for each the « best stationary value » by minimizing the empirical loss

$$R(\hat{y}) = \frac{1}{n} \sum_{i=1}^n \ell(y_i, \hat{y})$$

$$R'(\hat{y}) = \frac{1}{n} \sum_{i=1}^n (-\alpha e^{\alpha(y_i - \hat{y})} + \alpha) = 0$$

$$\Rightarrow \frac{1}{n} \sum_{i=1}^n e^{\alpha(y_i - \hat{y})} = 1$$

$$\Rightarrow \frac{1}{n} \sum_{i=1}^n e^{\alpha y_i} = e^{\hat{y}}$$

$$\Rightarrow \hat{y} = \log \frac{1}{n} \sum_{i=1}^n e^{\alpha y_i} = \text{softmax}(\alpha Y) - \log n$$



# 1<sup>st</sup> Approach – simple approach

- Advantages:
  - Simple model
  - Explainable
  - Use the real loss function LinEx
- Disadvantages:
  - Too many parameters (about 10000 lines in predict\_table)
  - Many of these parameters are correlated
  - Might have large variance
- Result:
  - 2.55 on the leader board

# 2<sup>nd</sup> approach – generalized version

- Our 1<sup>st</sup> approach can be regarded as a **linear regression** model
  - where the feature vector is “one-hot” encoding for all possible (ASS\_ASSIGNMENT, slot, dayofweek) tuples.
- We extend it by considering more features (month, day\_off)
- More general: consider **all possible combinations of all features!**
- => Linear LinEx Regression on Combined Features

# Combined Feature matrix

- For the row (ASS\_ASSIGNMENT, dayofweek, month, slot) = (Crises, 5, 1, 0)
- We associate a vector of 0 and 1's, where the 1's are in columns corresponding to

# Example for (ASS\_ASSIGNMENT, dayofweek, month, slot) = (Crises,5,1,0)

- (ASS\_ASSIGNMENT=Crises, dayofweek=5, month=1, slot=0)
- (dayofweek=5, month=1, slot=0)
- (ASS\_ASSIGNMENT=Crises, month=1, slot=0)
- (ASS\_ASSIGNMENT=Crises, dayofweek=5, slot=0)
- (ASS\_ASSIGNMENT=Crises, dayofweek=5, month=1)
- (month=1, slot=0)
- (dayofweek=5, slot=0)
- (dayofweek=5, month=1)
- (ASS\_ASSIGNMENT=Crises, slot=0)
- (ASS\_ASSIGNMENT=Crises, month=1)
- (ASS\_ASSIGNMENT=Crises, dayofweek=5)
- (slot=0)
- (month=1)
- (dayofweek=5)
- (ASS\_ASSIGNMENT=Crises)
- () (intercept)

A feature matrix of shape (1030829,147784)!

But each row has only **16 non-zero terms**

=> We use `scipy.sparse.csr_matrix` s

# Linear linex regression

Loss function

$$\frac{1}{n} \sum_{i=1}^n \ell(y_i, x_i^\top \theta) + \frac{\lambda}{2} \|\theta\|_2^2$$

where

$$\ell(x, y) = \text{LinEx}(x, y) = \exp(\alpha(x - y)) - \alpha(x - y) - 1$$

# Learning algorithm

- Stochastic gradient descent with variance reduction
- SVRG (Stochastic Variance Reduced Gradient) algorithm

# SVRG

**Input:** starting point  $\theta_0$ , learning rate  $\eta > 0$

Put  $\tilde{\theta}^1 \leftarrow \theta$

For  $k = 1, 2, \dots$  until convergence do

1. Put  $\theta_0^k \leftarrow \tilde{\theta}^1$

2. Compute  $\mu = \nabla f(\tilde{\theta}^k)$

3. For  $t = 0, \dots, m - 1$ :

- Pick uniformly at random  $i$  in  $\{1, \dots, n\}$
- Apply the step

$$\theta_{t+1}^k \leftarrow \theta_t^k - \eta(\nabla f_i(\theta_t^k) - \nabla f_i(\tilde{\theta}^k) + \mu)$$

Set

$$\tilde{\theta}^k \leftarrow \frac{1}{m} \sum_{t=1}^m \theta_t^k$$

**Return** last  $\theta_t^k$

# 2<sup>nd</sup> approach – generalized version

- Advantages:
  - The loss function is convex
  - Very general, containing many possible approaches as special case
  - Explainable
  - Use the real loss function LinEx
- Disadvantages:
  - Many parameters (about 150000 of them)
  - Hard to optimize
- Result:
  - 1.99 on the leader board



# 3<sup>rd</sup> approach – Random Forest

- Use all the features in feature engineering
- No categorical values in sklearn -> one-hot encoding
- Remove Evenements and Gestion Amex
- Cross validation (80% training, 20% testing)
- Multiply by  $C = 2.4$

# 3<sup>rd</sup> approach – Random Forest

- Advantages:
  - Robust model
  - Existing library
  - Relatively good results
- Disadvantages:
  - Parameter tuning
  - Hard to use a custom loss function
- Result:
  - 1.175 on the leaderboard

# Conclusion

- For prediction, when collecting data on past time, make sure this data will also be available for future times, otherwise they are not useful features for prediction
- A lot of features can be created on DATE and it can be enough when the data actually mostly depends on DATE